

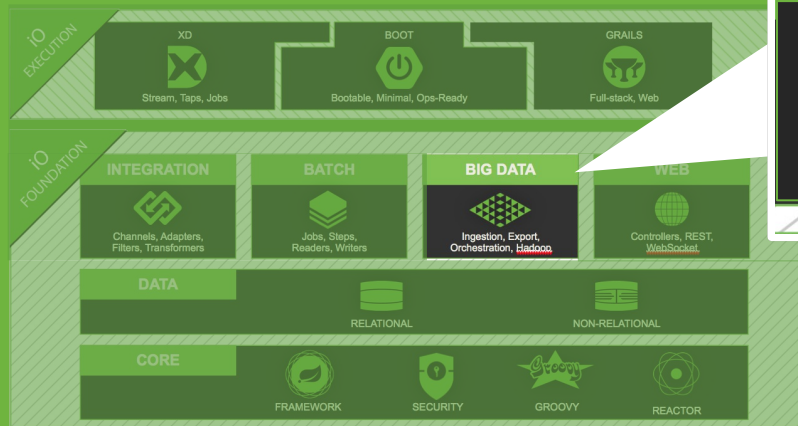


Getting started with Spring Data and Apache Hadoop

Thomas Risberg
Janne Valkealahti

Spring IO FOUNDATION:

Spring for Apache Hadoop



Just a Note




in the making of this presentation

About us ...

- Thomas
 - Working on the Spring Data engineering team at Pivotal
 - Joined Spring Framework team in 2003 working on JDBC support
 - co-author of “Professional Java Development with Spring Framework” from Wrox 2005 and “Spring Data” book from O'Reilly 2012
- Janne
 - Member of the Spring Data engineering team at Pivotal
 - contributes to Spring for Apache Hadoop and Spring XD projects
 - Previously Consultant for SpringSource vFabric team
 - 10-year career at a biggest online stock broker in Finland

About Apache Hadoop

- An Apache Project
- Modeled after Google File System and Map Reduce
- Provides:
 - Distributed file system
 - Map Reduce
 - General resource management for workloads with YARN (Hadoop v2)
 - Started as open source project at Yahoo and Facebook
 - Initial development by Doug Cutting and Mike Cafarella
- We hope you attended “Hadoop - Just the Basics for Big Data Rookies” with Adam Shook earlier; we will not cover Hadoop itself in detail today



About Spring Data

- Bring classic Spring value propositions to new data technologies
 - Productivity
 - Programming model consistency
- Support for new data technologies like NOSQL databases, Hadoop, SOLR, ElasticSearch, Querydsl
- Many entry points to use
 - Low level data access and Opinionated APIs
 - Repository Support and Object Mapping
 - Guidance



Hadoop trends

- Many organizations are currently using or evaluating Hadoop
- One common usage is Hadoop HDFS as a “data-lake” landing zone
 - *Collect all data and store it in HDFS, worry about analysis later*
- Many companies are now looking to YARN for running non-map-reduce workloads on a Hadoop cluster
- Lots of interest in using SQL on top of HDFS data – Hive/Stinger, Impala and HAWQ



What we will be talking about today

- Getting started with:
 - running **Apache Hadoop** for development/testing
 - writing map-reduce jobs for **Apache Hadoop**
 - writing apps using **Spring for Apache Hadoop**
 - writing apps with **Spring Yarn**

Getting Started with Hadoop

... can seem like an uphill struggle at times



... let's take one step at a time

Some ways to get started

1. Standalone Mode
2. Pre-configured VM
3. Pseudo-distributed cluster


Hadoop in Standalone Mode

- Download Apache Hadoop from
 - <http://hadoop.apache.org/releases.html#Download>
- Create a directory and unzip the download, set PATH and test

```
~$ mkdir ~/test
~$ cd ~/test
~/test$ tar xvzf ~/Downloads/hadoop-1.2.1-bin.tar.gz
~/test$ export HADOOP_INSTALL=~/test/hadoop-1.2.1
~/test$ export PATH=$PATH:$HADOOP_INSTALL/bin:$HADOOP_INSTALL/sbin
~/test$ export JAVA_HOME=/usr/lib/jvm/java-6-openjdk-amd64
~/test$ hadoop version
Hadoop 1.2.1
Subversion https://svn.apache.org/repos/asf/hadoop/common/branches/branch-1.2 -r 1503152
Compiled by mattf on Mon Jul 22 15:23:09 PDT 2013
From source with checksum 6923c86528809c4e7e6f493b6b413a9a
This command was run using /home/trisberg/test/hadoop-1.2.1/hadoop-core-1.2.1.jar
```

Our first Map Reduce job – TweetHashtagCounter ...


- We will count the number of occurrences of #hashtags in a collection of tweets collected during the 2013 NBA Finals
- Based on “Word Count” example from “MapReduce Design Patterns” book
- We need
 - a Mapper class
 - and a Reducer class
 - and a driver class



Some input data – tweets captured during NBA finals

```
{
  "id": 348115421360164864,
  "text": "RT @NBA: The Best of the 2013 #NBAFinals set to 'Radioactive' by Imagine Dragons! http://t.co/EA198meYpC",
  "createdAt": 1371832158000,
  "fromUser": "I_Nikki_I",
  ...
  "retweetedStatus": {
    "id": 348111916452950016,
    "text": "The Best of the 2013 #NBAFinals set to 'Radioactive' by Imagine Dragons! http://t.co/EA198meYpC",
    "createdAt": 1371831323000,
    "fromUser": "NBA",
    ...
  },
  ...
  "entities": {
    "hashTags": [{
      "text": "NBAFinals",
      "indices": [30, 40]
    }]
  },
  "retweet": true
}
```

The data file has the entire JSON document for each tweet on a single line



Our first Mapper class

```
public class TweetCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
  
    private final static IntWritable ONE = new IntWritable(1);  
    private final ObjectMapper mapper = new ObjectMapper(new JsonFactory());  
  
    @Override  
    protected void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
  
        Map<String, Object> tweet = mapper.readValue(value.toString(),  
            new TypeReference<HashMap<String, Object>>(){});  
        Map<String, Object> entities = (Map<String, Object>) tweet.get("entities");  
        List<Map<String, Object>> hashTagEntries = null;  
        if (entities != null) {  
            hashTagEntries = (List<Map<String, Object>>) entities.get("hashTags");  
        }  
        if (hashTagEntries != null && hashTagEntries.size() > 0) {  
            for (Map<String, Object> hashTagEntry : hashTagEntries) {  
                String hashTag = hashTagEntry.get("text").toString();  
                context.write(new Text(hashTag), ONE);  
            }  
        }  
    }  
}
```

Our first Reducer class

```
public class IntSumReducer extends Reducer<Text, IntWritable, Text, IntWritable>{

    @Override
    protected void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable value : values) {
            sum += value.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

Our first Driver class

```
public class TweetHashTagCounter {  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        String[] myArgs = new GenericOptionsParser(conf, args).getRemainingArgs();  
        if (myArgs.length != 2) {  
            System.err.println("Usage: TweetHashTagCounter <input path> <output path>");  
            System.exit(-1);  
        }  
        Job job = Job.getInstance(conf, "Tweet Hash Tag Counter");  
        job.setJarByClass(TweetHashTagCounter.class);  
  
        FileInputFormat.addInputPath(job, new Path(myArgs[0]));  
        FileOutputFormat.setOutputPath(job, new Path(myArgs[1]));  
  
        job.setMapperClass(TweetCountMapper.class);  
        job.setReducerClass(IntSumReducer.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Need to build the app – our pom.xml

```
<project xmlns="http://maven.apache.org/POM/4.0.0"
...
  <groupId>com.springdeveloper.hadoop</groupId>
  <artifactId>tweet-counts-hadoop</artifactId>
  <version>0.1.0</version>
  <packaging>jar</packaging>
  <name>Tweet Counts</name>
...
  <properties>
    <hadoop.version>1.2.1</hadoop.version>
  </properties>
...
  <dependencies>
    <dependency>
      <groupId>org.apache.hadoop</groupId>
      <artifactId>hadoop-core</artifactId>
      <version>${hadoop.version}</version>
    </dependency>
...
</dependencies>
```


Example code repository

- All code for the Hadoop HDFS and Map Reduce examples can be downloaded from GitHub
- Repository
 - <https://github.com/trisberg/springone-hadoop.git>

```
$ cd ~  
$ git clone https://github.com/trisberg/springone-hadoop.git  
$ cd ~/springone-hadoop
```

Let's build and run the app

```
$ cd ~/springone-hadoop/tweet-counts-hadoop
$ export HADOOP_INSTALL=~/.test/hadoop-1.2.1
$ export PATH=$PATH:$HADOOP_INSTALL/bin:$HADOOP_INSTALL/sbin
$ export JAVA_HOME=/usr/lib/jvm/java-6-openjdk-amd64
$ mvn clean install

...
[INFO] --- maven-jar-plugin:2.3.1:jar (default-jar) @ tweet-counts-hadoop ---
[INFO] Building jar: /home/trisberg/springone-hadoop/tweet-counts-hadoop/target/tweet-counts-hadoop-0.1.0.jar
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 7.024s
...
$ export HADOOP_CLASSPATH=~/.springone-hadoop/tweet-counts-hadoop/target/tweet-counts-hadoop-0.1.0.jar
$
$ hadoop com.springdeveloper.hadoop.TweetHashTagCounter ~/.springone-hadoop/data/nbatweets-small.txt ~/.springone-hadoop/output
```

App log output ...

```
13/09/01 13:24:04 INFO mapred.LocalJobRunner:
13/09/01 13:24:04 INFO mapred.Task: Task attempt_local868926382_0001_r_000000_0 is allowed
to commit now
13/09/01 13:24:04 INFO output.FileOutputCommitter: Saved output of task
'attempt_local868926382_0001_r_000000_0' to /home/trisberg/springone-hadoop/output
13/09/01 13:24:04 INFO mapred.LocalJobRunner: reduce > reduce
13/09/01 13:24:04 INFO mapred.Task: Task 'attempt_local868926382_0001_r_000000_0' done.
13/09/01 13:24:04 INFO mapred.JobClient: map 100% reduce 100%
13/09/01 13:24:04 INFO mapred.JobClient: Job complete: job_local868926382_0001
13/09/01 13:24:04 INFO mapred.JobClient: Counters: 17
13/09/01 13:24:04 INFO mapred.JobClient:   File Output Format Counters
13/09/01 13:24:04 INFO mapred.JobClient:     Bytes Written=9894
13/09/01 13:24:04 INFO mapred.JobClient:   File Input Format Counters
13/09/01 13:24:04 INFO mapred.JobClient:     Bytes Read=14766958
13/09/01 13:24:04 INFO mapred.JobClient:   FileSystemCounters
13/09/01 13:24:04 INFO mapred.JobClient:     FILE_BYTES_READ=29581326
13/09/01 13:24:04 INFO mapred.JobClient:     FILE_BYTES_WRITTEN=193326
...
13/09/01 13:24:04 INFO mapred.JobClient:   Reduce output records=836
13/09/01 13:24:04 INFO mapred.JobClient:   Map output records=2414
```

And the results ...

```
$ more ~/springone-hadoop/output/part-r-00000 | grep NBA
2013NBACHamps 1
NBA 474
NBA2K12 1
NBA2K14 2
NBA4ARAB 1
NBAAllStar 2
NBACHampions 3
NBACHamps 4
NBADraft 3
NBAFINALS 8
NBAFinals 88
NBAFrance 1
NBALIVE14 1
NBARigged 1
NBASoutheast 2
NBATV 2
...
```

Developer observations on Hadoop

- For Spring developers, Hadoop has a fairly poor out-of-the-box programming model
- Lots of low-level configuration and exception handling code
- Non trivial applications often become a collection of scripts calling Hadoop command line applications
- Spring aims to simplify development for Hadoop applications
 - Leverage Spring's configuration features in addition to several Spring eco-system projects

Spring for Apache Hadoop



*Spring for Apache Hadoop
provides extensions to Spring,
Spring Batch, and Spring
Integration to build manageable
and robust pipeline solutions
around Hadoop.*

Spring for Apache Hadoop – Features

- Consistent programming and declarative configuration model
 - Create, configure, and parametrize Hadoop connectivity and all job types
 - Environment profiles – easily move application from dev to qa to production
- Developer productivity
 - Create well-formed applications, not spaghetti script applications
 - Simplify HDFS access and FsShell API with support for JVM scripting
 - Runner classes for MR/Pig/Hive/Cascading for small workflows
 - Helper “Template” classes for Pig/Hive/HBase

Spring for Apache Hadoop – Use Cases

- Apply across a wide range of use cases
 - Ingest: Events/JDBC/NoSQL/Files to HDFS
 - Orchestrate: Hadoop Jobs
 - Export: HDFS to JDBC/NoSQL
- Spring Integration and Spring Batch make this possible

Spring for Apache Hadoop – Status

- 1.0 GA in February 2013 – supported up to Hadoop 1.0.4
- 1.0.1 GA last week – supports all Hadoop 1.x stable, 2.0-alpha and 2.1-beta
- Default is Apache Hadoop 1.2.1 stable
- Distribution specific “flavors” via a suffix on version:
 - 1.0.1.RELEASE-cdh4 Cloudera CDH 4.3.1
 - 1.0.1.RELEASE-hdp13 Hortonworks HDP 1.3
 - 1.0.1.RELEASE-phd1 Pivotal HD 1.0
 - 1.0.1.RELEASE-hadoop21 Hadoop 2.1.0-beta

Spring for Apache Hadoop – Future

- New structure for 2.0
 - New sub-projects
 - Core M/R, FSShell, Hive, Pig etc. and basic configuration
 - Batch is separate with separate namespace
 - Cascading separate with separate namespace
 - Test sub-project for integration testing
 - Adding **spring-yarn** sub-project for 2.0 based builds
 - Just released first 2.0.0.M1 milestone release

Core and Batch

- Running Map Reduce jobs
- HDFS shell scripting
- Running Pig and Hive scripts
- Configuration
- Configuring batch jobs with Spring Batch

Examples built using spring-data-hadoop 1.0.1.RELEASE

Hadoop Configuring M/R

- Standard Hadoop APIs

```
Configuration conf = new Configuration();  
Job job = Job.getInstance(conf, "Tweet Hash Tag Counter");  
job.setJarByClass(TweetHashTagCounter.class);
```

```
FileInputFormat.addInputPath(job, new Path(myArgs[0]));  
FileOutputFormat.setOutputPath(job, new Path(myArgs[1]));
```

```
job.setMapperClass(TweetCountMapper.class);  
job.setReducerClass(IntSumReducer.class);  
job.setOutputKeyClass(Text.class);  
job.setOutputValueClass(IntWritable.class);
```

```
System.exit(job.waitForCompletion(true) ? 0 : 1);
```

Configuring Hadoop with Spring

```
<context:property-placeholder location="hadoop-dev.properties"/>
```

applicationContext.xml

```
<hdp:configuration>
```

```
  fs.default.name=${hd.fs}
```

```
  mapred.job.tracker=${hd.jt}
```

```
</hdp:configuration>
```

```
<hdp:job id="word-count-job"
```

```
  input-path="${input.path}"
```

```
  output-path="${output.path}"
```

```
  jar="hadoop-examples.jar"
```

```
  mapper="examples.WordCount.WordMapper"
```

```
  reducer="examples.WordCount.IntSumReducer"/>
```

Automatically determines
Output key and class

```
<hdp:job-runner id="runner" job-ref="word-count-job"
```

```
  run-at-startup="true" />
```

```
input.path=/wc/input/
```

```
output.path=/wc/word/
```

```
hd.fs=hdfs://localhost:8020
```

```
hd.jt=localhost:8021
```

hadoop-dev.properties

Injecting Jobs

- Use DI to obtain reference to Hadoop Job
 - Perform additional runtime configuration and submit

```
public class WordService {  
  
    @Autowired  
    private Job mapReduceJob;  
  
    public void processWords() {  
        mapReduceJob.submit();  
    }  
}
```

Streaming Jobs and Environment Configuration

```
bin/hadoop jar hadoop-streaming.jar \  
  -input /wc/input -output /wc/output \  
  -mapper /bin/cat -reducer /bin/wc \  
  -files stopwords.txt
```

```
<context:property-placeholder location="hadoop-dev.properties"/>  
  
<hdp:streaming id="wc" input-path="${input}" output-path="${output}"  
  mapper="${cat}" reducer="${wc}"  
  files="classpath:stopwords.txt">  
</hdp:streaming>
```

hadoop-dev.properties

```
input.path=/wc/input/  
output.path=/wc/word/  
hd.fs=hdfs://localhost:9000
```

hadoop-qa.properties

```
input.path=/gutemberg/input/  
output.path=/gutemberg/word/  
hd.fs=hdfs://darwin:9000
```

```
Java -Denv=dev -jar SpringLauncher.jar applicationContext.xml
```

HDFS and Hadoop Shell as APIs

- Access all “bin/hadoop fs” commands through Spring’s FsShell helper class
 - mkdir, chmod, test

```
class MyScript {  
    @Autowired FsShell fsh;  
  
    @PostConstruct void init() {  
        String outputDir = "/data/output";  
        if (fsShell.test(outputDir)) {  
            fsShell.rmr(outputDir);  
        }  
    }  
}
```


HDFS and Hadoop Shell as APIs



```
copy-files.groovy
```

```
// use the shell (made available under variable fsh)

if (!fsh.test(inputDir)) {
    fsh.mkdir(inputDir);
    fsh.copyFromLocal(sourceFile, inputDir);
    fsh.chmod(700, inputDir)
}
if (fsh.test(outputDir)) {
    fsh.rmr(outputDir)
}
```

HDFS and Hadoop Shell as APIs



- Reference script and supply variables in application configuration

app-context.xml

```
<script id="setupScript" location="copy-files.groovy">
  <property name="inputDir" value="${wordcount.input.path}"/>
  <property name="outputDir" value="${wordcount.output.path}"/>
  <property name="sourceFile" value="${localSourceFile}"/>
</script>
```

Small workflows

- Often need the following steps
 - Execute HDFS operations before job
 - Run MapReduce Job
 - Execute HDFS operations after job completes
- Spring's JobRunner helper class sequences these steps
 - Can reference multiple scripts with comma delimited names

```
<hdp:job-runner id="runner" run-at-startup="true"  
    pre-action="setupScript"  
    job="wordcountJob"  
    post-action="tearDownScript"/>
```

Runner classes

- Similar runner classes available for Hive and Pig
- Implement JDK callable interface
- Easy to schedule for simple needs using Spring

```
<hdp:job-runner id="runner" run-at-startup="false"
                pre-action="setupScript"
                job="wordcountJob"
                post-action="tearDownScript"/>


<task:scheduled-tasks>
  <task:scheduled ref="runner" method="call" cron="3/30 * * * * ?"/>
</task:scheduled-tasks>
```

- Can later 'graduate' to use Spring Batch for more complex workflows
 - Start simple and grow, reusing existing configuration

Our first Spring Configured Map Reduce job

- We will reuse the TweetHashTagCounter example
- Loosely based on “Spring Word Count” example from “Spring Data” book
- We need an application context
 - and a properties file
 - and a driver class

<https://github.com/trisberg/springone-hadoop.git>



Our application context

```
<context:property-placeholder location="hadoop.properties"/>
```

```
<configuration>
```

```
    fs.default.name=${hd.fs}
```

```
    mapred.job.tracker=${hd.jt}
```

```
</configuration>
```

```
<job id="tweetCountJob"
```

```
    input-path="${tweetcount.input.path}"
```

```
    output-path="${tweetcount.output.path}"
```

```
    libs="file:${app.repo}/tweet-counts-hadoop-0.1.0.jar"
```

```
    mapper="com.springdeveloper.hadoop.TweetCountMapper"
```

```
    reducer="com.springdeveloper.hadoop.IntSumReducer"/>
```

```
<script id="setupScript" location="file-prep.groovy">
```

```
    <property name="localSourceFile" value="${app.home}/${localSourceFile}"/>
```

```
    <property name="inputDir" value="${tweetcount.input.path}"/>
```

```
    <property name="outputDir" value="${tweetcount.output.path}"/>
```

```
</script>
```

```
<job-runner id="runner" run-at-startup="true"
```

```
    pre-action="setupScript"
```

```
    job-ref="tweetCountJob" />
```

```
hd.fs=hdfs://sandbox:8020
```

```
hd.jt=sandbox:50300
```

```
tweetcount.input.path=/tweets/input
```

```
tweetcount.output.path=/tweets/results
```

```
localSourceFile=data/nbatweets-small.txt
```

Our Spring Driver class

```
public class TweetCount {  
  
    private static final Log log = LoggerFactory.getLog(TweetCount.class);  
  
    public static void main(String[] args) throws Exception {  
        AbstractApplicationContext context = new ClassPathXmlApplicationContext(  
            "/META-INF/spring/application-context.xml", TweetCount.class);  
        log.info("TweetCount Application Running");  
        context.registerShutdownHook();  
    }  
}
```

A pom.xml to build and run the app – part 1

```
<project xmlns="http://maven.apache.org/POM/4.0.0"
...
  <properties>
    <spring.framework.version>3.2.4.RELEASE</spring.framework.version>
    <spring.hadoop.version>1.0.1.RELEASE</spring.hadoop.version>
  </properties>
  <dependencies>
    <dependency>
      <groupId>com.springdeveloper.hadoop</groupId>
      <artifactId>tweet-counts-hadoop</artifactId>
      <version>0.1.0</version>
      <scope>runtime</scope>
    </dependency>
    <dependency>
      <groupId>org.springframework</groupId>
      <artifactId>spring-context-support</artifactId>
      <version>${spring.framework.version}</version>
    </dependency>
    <dependency>
      <groupId>org.springframework.data</groupId>
      <artifactId>spring-data-hadoop</artifactId>
      <version>${spring.hadoop.version}</version>
    </dependency>
  </dependencies>
...
```


A pom.xml to build and run the app – part 2

```
...
    <dependency>
        <groupId>org.codehaus.groovy</groupId>
        <artifactId>groovy</artifactId>
        <version>1.8.5</version>
        <scope>runtime</scope>
    </dependency>
    <dependency>
        <groupId>log4j</groupId>
        <artifactId>log4j</artifactId>
        <version>1.2.14</version>
    </dependency>
...
</dependencies>
<plugins>
    <plugin>
        <groupId>org.codehaus.mojo</groupId>
        <artifactId>appassembler-maven-plugin</artifactId>
        <version>1.2.2</version>
    </plugin>
...
</plugins>
```

Testing with Hadoop – using a pre-configured VM

- VMs “ready to run” - most distro companies provide one:
 - Hortonworks Sandbox HDP 1.3 and HDP 2.0
 - Pivotal HD 1.0 Single Node VM
 - Cloudera Quickstart CDH4
- Which one to use? Depends on what your company uses.
- If starting from scratch Hortonworks Sandbox HDP 1.3 is based on Hadoop 1.2.0 and a good place to start ...
 - HDFS configured to listen on the VM network making it easy to connect from host system
 - Uses only 2GB of memory making it easy to use on laptops
 - Compatible with Spring for Apache Hadoop 1.0.1.RELEASE and its transitive dependencies

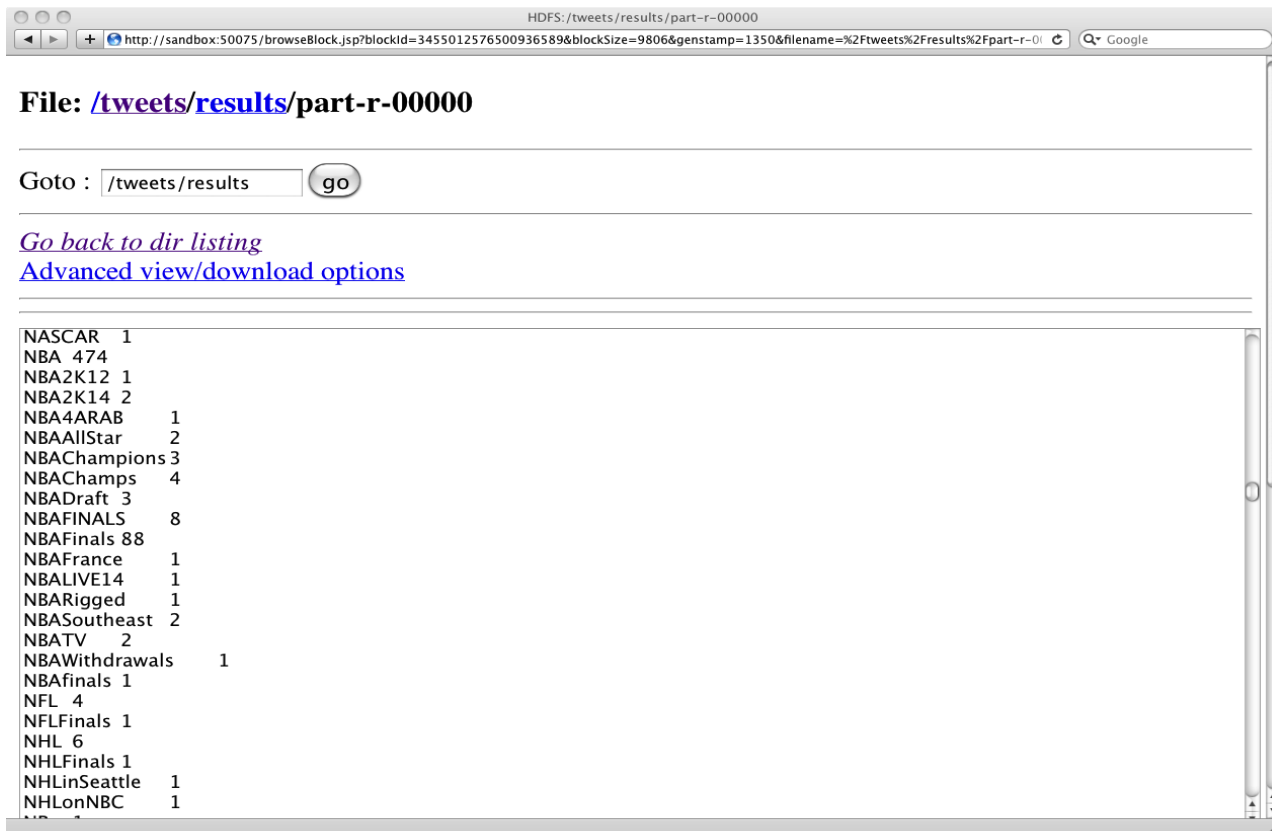
Let's build and run the Spring app

```
$ cd ~/springone-hadoop/tweet-counts-spring
$ mvn clean package
...
[INFO] --- maven-antrun-plugin:1.3:run (config) @ tweet-counts-spring ---
[INFO] Executing tasks
    [copy] Copying 1 file to /home/trisberg/springone-hadoop/tweet-counts-
spring/target/appassembler/data
[INFO] Executed tasks
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 11.710s
...
$ sh ./target/appassembler/bin/tweetcount
```

App log output ...

```
...
13/09/01 16:28:42 INFO mapreduce.JobRunner: Starting job [tweetCountJob]
13/09/01 16:28:42 WARN mapred.JobClient: No job jar file set. User classes may not be
found. See JobConf(Class) or JobConf#setJar(String).
13/09/01 16:28:42 INFO input.FileInputFormat: Total input paths to process : 1
13/09/01 16:28:43 INFO mapred.JobClient: Running job: job_201308311801_0002
13/09/01 16:28:44 INFO mapred.JobClient: map 0% reduce 0%
13/09/01 16:29:03 INFO mapred.JobClient: map 25% reduce 0%
13/09/01 16:29:06 INFO mapred.JobClient: map 78% reduce 0%
13/09/01 16:29:08 INFO mapred.JobClient: map 100% reduce 0%
13/09/01 16:29:20 INFO mapred.JobClient: map 100% reduce 33%
13/09/01 16:29:23 INFO mapred.JobClient: map 100% reduce 100%
13/09/01 16:29:25 INFO mapred.JobClient: Job complete: job_201308311801_0002
...
13/09/01 16:29:25 INFO mapred.JobClient: Reduce input records=2414
13/09/01 16:29:25 INFO mapred.JobClient: Reduce input groups=836
...
13/09/01 16:29:25 INFO mapred.JobClient: Map output records=2414
13/09/01 16:29:25 INFO mapreduce.JobRunner: Completed job [tweetCountJob]
...
```

And the results ...



The screenshot shows a web browser window with the address bar displaying a URL from a sandbox environment. The page title is "HDFS: /tweets/results/part-r-00000". Below the title, the file path is shown as "File: /tweets/results/part-r-00000". There is a "Goto" field with the text "/tweets/results" and a "go" button. Below this, there are two links: "Go back to dir listing" and "Advanced view/download options". The main content area displays a list of sports-related terms and their corresponding counts, such as "NASCAR 1", "NBA 474", "NBA2K12 1", etc.

File: [/tweets/results/part-r-00000](#)

Goto :

[Go back to dir listing](#)
[Advanced view/download options](#)

NASCAR	1
NBA	474
NBA2K12	1
NBA2K14	2
NBA4ARAB	1
NBAAllStar	2
NBACHampions	3
NBACHamps	4
NBADraft	3
NBAFINALS	8
NBAFinals 88	
NBAFrance	1
NBALIVE14	1
NBARigged	1
NBASoutheast	2
NBATV	2
NBAWithdrawals	1
NBAfinals	1
NFL	4
NFLFinals	1
NHL	6
NHLFinals	1
NHLinSeattle	1
NHLonNBC	1

Spring's PigRunner



- Execute a small Pig workflow

```
<pig-factory job-name="analysis" properties-location="pig-server.properties"/>

<script id="hdfsScript" location="copy-files.groovy">
  <property name="sourceFile" value="${localSourceFile}"/>
  <property name="inputDir" value="${inputDir}"/>
  <property name="outputDir" value="${outputDir}"/>
</script>

<pig-runner id="pigRunner" pre-action="hdfsScript" run-at-startup="true">
  <script location="wordCount.pig">
    <arguments>
      inputDir=${inputDir}
      outputDir=${outputDir}
    </arguments>
  </script>
</pig-runner>
```

PigTemplate - Configuration

- Helper class that simplifies the programmatic use of Pig
 - Common tasks are one-liners

```
<pig-factory id="pigFactory" properties-location="pig-server.properties"/>
```

```
<pig-template pig-factory-ref="pigFactory"/>
```

- Similar XxxTemplate helper classes for Hive and HBase

PigTemplate – Programmatic Use

```
public class PigPasswordRepository implements PasswordRepository {
    @Autowired
    private PigTemplate pigTemplate;
    @Autowired
    private String outputDir;
    private String pigScript = "classpath:password-analysis.pig";

    public void processPasswordFile(String inputFile) {

        Properties scriptParameters = new Properties();
        scriptParameters.put("inputDir", inputFile);
        scriptParameters.put("outputDir", outputDir);

        pigTemplate.executeScript(pigScript, scriptParameters);
    }
}
```



Pig example using Spring

- We will use the output from the TweetHashTagCounter example
- Sort and select the top 10 #hashtags
- We need an application context
 - With an embedded Pig server
 - and a properties file
 - and a driver class



Our Pig script

```
hashtags = LOAD '$inputDir' USING PigStorage('\t') AS (hashtag:chararray, count:int);
sorted = ORDER hashtags BY count DESC;
top10 = LIMIT sorted 10;
STORE top10 INTO '$outputDir';
```



DEMO - Pig

```
<context:property-placeholder location="hadoop.properties,pig.properties"/>

<configuration>
    fs.default.name=${hd.fs}
    mapred.job.tracker=${hd.jt}
</configuration>

<script id="hdfsScript" language="groovy" location="file-prep.groovy">
    <property name="outputDir" value="${pig.outputPath}"/>
</script>

<pig-factory exec-type="MAPREDUCE" properties-location="pig-server.properties"/>

<pig-runner id="pigRunner"
    pre-action="hdfsScript"
    run-at-startup="true" >
    <script location="tweet-analysis.pig">
        <arguments>
            inputDir=${pig.inputPath}
            outputDir=${pig.outputPath}
        </arguments>
    </script>
</pig-runner>
```

<https://github.com/trisberg/springone-hadoop.git>

Hive example using Spring


- We will count the number of retweets per original user account found in the collection of tweets collected during the 2013 NBA Finals
- Sort and select the top 10 users based on the number of retweets found – this should give us the influential users
- We need an application context
 - With and embedded Hive server
 - and a properties file
 - and a driver class



Same input data – tweets captured during NBA finals

```
{
  "id": 348115421360164864,
  "text": "RT @NBA: The Best of the 2013 #NBAFinals set to 'Radioactive' by Imagine Dragons! http://t.co/EA198meYpC",
  "createdAt": 1371832158000,
  "fromUser": "I_Nikki_I",
  ...
  "retweetedStatus": {
    "id": 348111916452950016,
    "text": "The Best of the 2013 #NBAFinals set to 'Radioactive' by Imagine Dragons! http://t.co/EA198meYpC",
    "createdAt": 1371831323000,
    "fromUser": "NBA",
    ...
  },
  ...
  "entities": {
    "hashTags": [{
      "text": "NBAFinals",
      "indices": [30, 40]
    }]
  },
  "retweet": true
}
```


The data file has the entire JSON document for each tweet on a single line



Our Hive script

```
create external table tweetdata (value STRING) LOCATION '/tweets/input';

select r.retweetedUser, count(r.retweetedUser) as count
  from tweetdata j
    lateral view json_tuple(j.value, 'retweet', 'retweetedStatus') t as retweet, retweetedStatus
    lateral view json_tuple(t.retweetedStatus, 'fromUser') r as retweetedUser
 where t.retweet = 'true'
group by r.retweetedUser order by count desc limit 10;
```



DEMO - Hive

```
<context:property-placeholder location="hive-jdbc.properties"/>


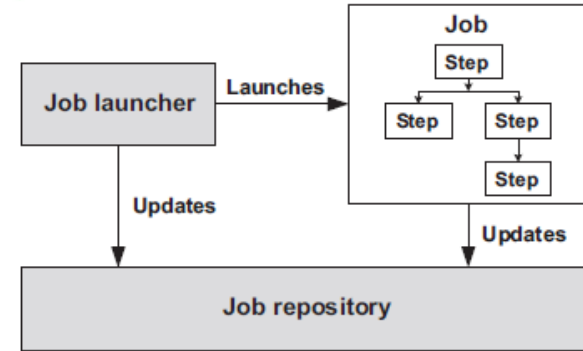
<bean id="dataSource" class="org.springframework.jdbc.datasource.SimpleDriverDataSource">
  <property name="driverClass" value="org.apache.hive.jdbc.HiveDriver"/>
  <property name="url" value="${hive.url}"/>
  <property name="username" value="${hive.user}"/>
</bean>

<bean id="jdbcTemplate" class="org.springframework.jdbc.core.JdbcTemplate">
  <constructor-arg ref="dataSource"/>
</bean>
```

<https://github.com/trisberg/springone-hadoop.git>

Spring Batch

- Framework for batch processing
 - Basis for JSR-352
- Born out of collaboration with Accenture in 2007
- Features
 - parsers, mappers, readers, writers
 - automatic retries after failure
 - periodic commits
 - synchronous and asynch processing
 - parallel processing
 - partial processing (skipping records)
 - non-sequential processing
 - job tracking and restart


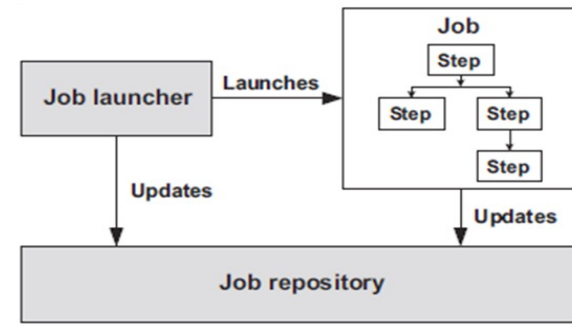


Spring Batch workflows for Hadoop

- Batch Ingest/Export
 - Examples
 - Read log files on local file system, transform and write to HDFS
 - Read from HDFS, transform and write to JDBC, HBase, MongoDB,...
- Batch Analytics
 - Orchestrate Hadoop based workflows with Spring Batch
 - Also orchestrate non-hadoop based workflows

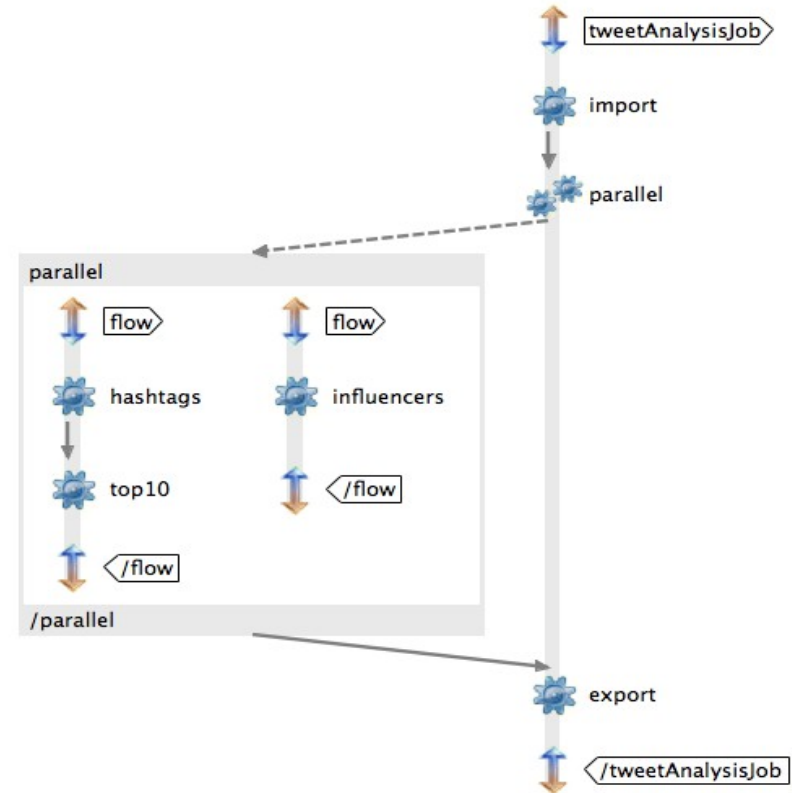
Hadoop Analytical workflow managed by Spring Batch

- Reuse same Batch infrastructure and knowledge to manage Hadoop workflows
- Step can be any Hadoop job type or HDFS script



Spring Batch Configuration for Hadoop

```
<batch:job id="tweetAnalysisJob">
  <batch:step id="import" next="parallel">
    <batch:tasklet ref="scriptTasklet"/>
  </batch:step>
  <batch:split id="parallel" task-executor="taskExec">
    <batch:flow>
      <batch:step id="hashtags" next="top10">
        <batch:tasklet ref="hashtag-tasklet" />
      </batch:step>
      <batch:step id="top10">
        <batch:tasklet ref="top10-tasklet" />
      </batch:step>
    </batch:flow>
    <batch:flow>
      <batch:step id="influencers">
        <batch:tasklet ref="influencers-tasklet" />
      </batch:step>
    </batch:flow>
  </batch:split>
  <batch:step id="export" parent="export-step"/>
</batch:job>
```

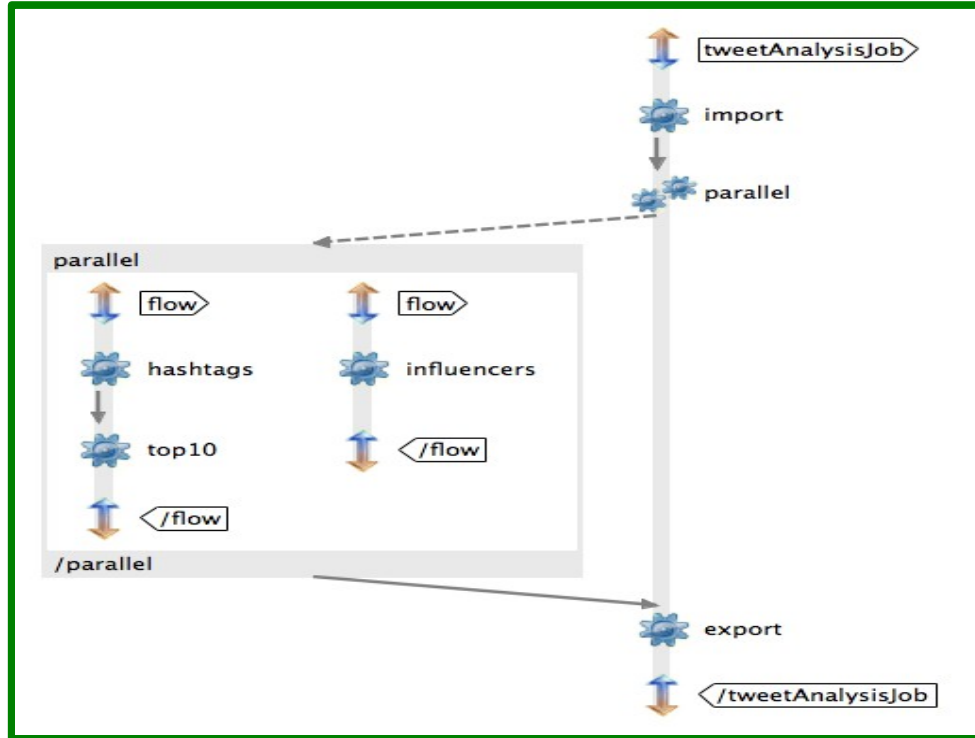


Exporting HDFS to JDBC

- Use Spring Batch's
 - MutliResourceItemReader + FlatFileItemReader
 - JdbcBatchItemWriter

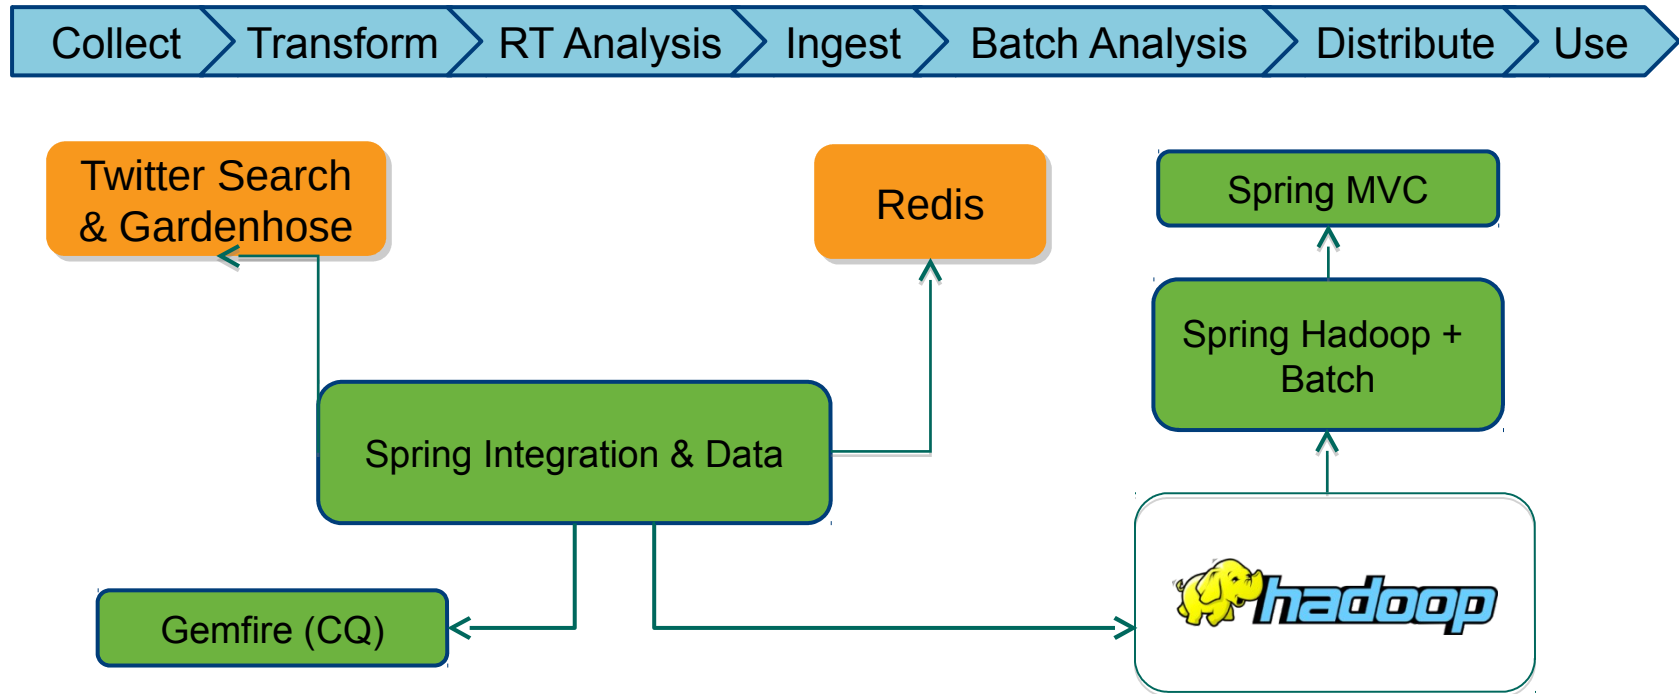
```
<step id="step1">
  <tasklet>
    <chunk reader="hdfsReader" processor="itemProcessor"
            writer="jdbcWriter"
            commit-interval="1000"
            skip-limit="3"/>
    </chunk>
  </tasklet>
</step>
```

DEMO - Batch

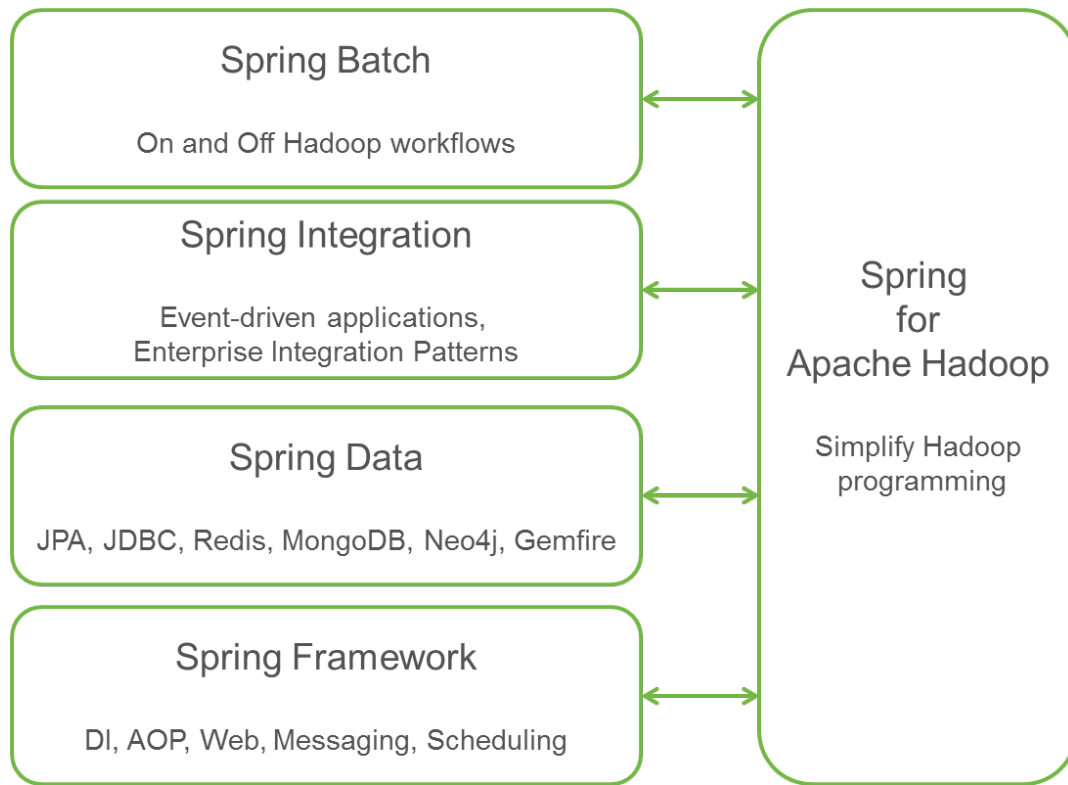


<https://github.com/trisberg/springone-hadoop.git>

Big Data problems are also integration problems



Relationship between Spring Projects



Next Steps – Spring XD

- New open source umbrella project to support common big data use cases
 - High throughput distributed data ingestion into HDFS
 - From a variety of input sources
 - Real-time analytics at ingestion time
 - Gathering metrics, counting values, Gemfire CQ...
 - On and off Hadoop workflow orchestration
 - High throughput data export
 - From HDFS to a RDBMS or NoSQL database.

Tackling Big Data Complexity with Spring

2:30 - 4:00 PM

SCCC Theatre

Don't miss!

Spring Yarn




Spring Yarn provides features from the Spring programming model to make developing Yarn applications as easy as developing regular Spring applications.

Hadoop Yarn

- Hadoop v1 vs. v2
- Is a Resource Scheduler
- Is not a Task Scheduler
- YARN != Hadoop v2
- MapReduce v2 is a YARN Application
- Big Investment – Re-use Outside of MapReduce

YARN Components



Spring Yarn

- Is a Framework
- Run Spring Contexts on YARN
- Application Configuration
- No Boilerplate for Something Simple
- Extend to Create more Complex Applications

Spring Yarn Concepts

- Configuration XML vs. JavaConfig (Milestone 2)
- Client
- Appmaster
- Container
- Bootstrap / Control

Concepts - Configuration

- Familiar Spring Config Styles
 - XML – Namespace

```
<beans ...>
```

```
  <yarn:configuration />
```

```
  <yarn:client />
```

```
  <yarn:appmaster />
```

```
  <yarn:container />
```

```
</beans>
```

Concepts – Configuration (Milestone 2)

- Familiar Spring Config Styles
 - JavaConfig – Builder / Configurers

```
@Configuration
@EnableYarn(enable=Enable.CLIENT)
class Config extends SpringYarnConfigurerAdapter
{
    @Override
    public void configure(YarnConfigBuilder config)
        throws ... {
        ...
    }
}
```

Concepts - Client

- Access Yarn Cluster
- Submit / Control Running Applications
- Launch Context for Appmaster
 - Config
 - Libraries (Localization)
 - Environment

Concepts - Appmaster

- Control the Running Application
- Appmaster is a main() of the Application
- Lifecycle
- Controls and Launches Containers
- Launch Context for Container
 - Config
 - Libraries (Localization
 - Environment

Concepts - Container

- Real Job or Task is Done Here
- Run / Do Something and Exit
- Interact with Custom Services

Concepts – Bootstrap / Control

- Application Context Having a YarnClient
 - Submit / Control
- CommandLineClientRunner
- Spring Boot
- Things to Remember
 - Dependencies for Hadoop Yarn Libs
 - Dependencies for Your Custom Code
 - Container Localized Files

Project Setup

- Custom Class Files / Context Configs
- Testing Files if Needed
- Spring Yarn Examples
- Normal Spring Project

```
src/main/java/.../MultiContextContainer.java  
src/main/resources/application-context.xml  
src/main/resources/appmaster-context.xml  
src/main/resources/container-context.xml  
src/test/java/.../MultiContextTests.java  
src/test/resources/MultiContextTests-context.xml
```

Demo

- Simple Example
 - Run Multiple Containers
 - Let Containers Just Exit
 - Application Master is Finished
 - Application is Completed

Testing with YARN

- Testing is Difficult
- Spring Yarn to Rescue
- Spring Test / Spring Yarn Test
- @MiniYarnCluster
- AbstractYarnClusterTests
- Yarn Configuration from a Mini Cluster

Test – Client Context Config

```
<beans ...>
```

```
<!-- where is our yarn config? -->
```

```
<yarn:localresources />
```

```
<yarn:environment />
```

```
<yarn:client app-name="myAppName">
```

```
  <yarn:master-runner />
```

```
</yarn:client>
```

```
</beans>
```

Test - JUnit

```
@ContextConfiguration
(loader=YarnDelegatingSmartContextLoader.class)
@MiniYarnCluster
public class AppTests extends AbstractYarnClusterTests {
    @Test
    public void testApp() throws IOException {
        YarnApplicationState state =
            SubmitApplicationAndWait();
        assertNotNull(state);
        assertTrue(state.equals(
            YarnApplicationState.FINISHED));
    }
}
```


Advanced Topic - Appmaster Services

- Link Between Appmaster and Container
 - Command / Control Container Internals
- Link Between Appmaster and Client
 - Command Your Custom Appmaster

Advanced Topic - Container Locality

- Task Accessing Data on HDFS
- Container “near” HDFS Blocks
 - On Nodes
 - On Racks

Advanced Topic - Spring Batch

- Execute Batch Partitioned Steps on Hadoop
- Proxy for Remote Job Repository
- Appmaster Runs the Batch Job

Spring Yarn Future?

- M2 planned for Q4
- Java Config support
- 2.1.x-beta Overhauls Yarn APIs
 - Incompatible with Hadoop 2.0 alpha based distributions
- Potential Extensions
 - Thrift
 - Heartbeating
 - Container Grid/Groups

Installing Hadoop



A couple of ways to install a small Hadoop cluster that can be used to test your new Hadoop applications.

Hortonworks HDP 1.3 Sandbox

- Download:
 - <http://hortonworks.com/products/hortonworks-sandbox/>
- VMs available for:
 - VirtualBox
 - VMware Fusion or Player
 - Hyper-V



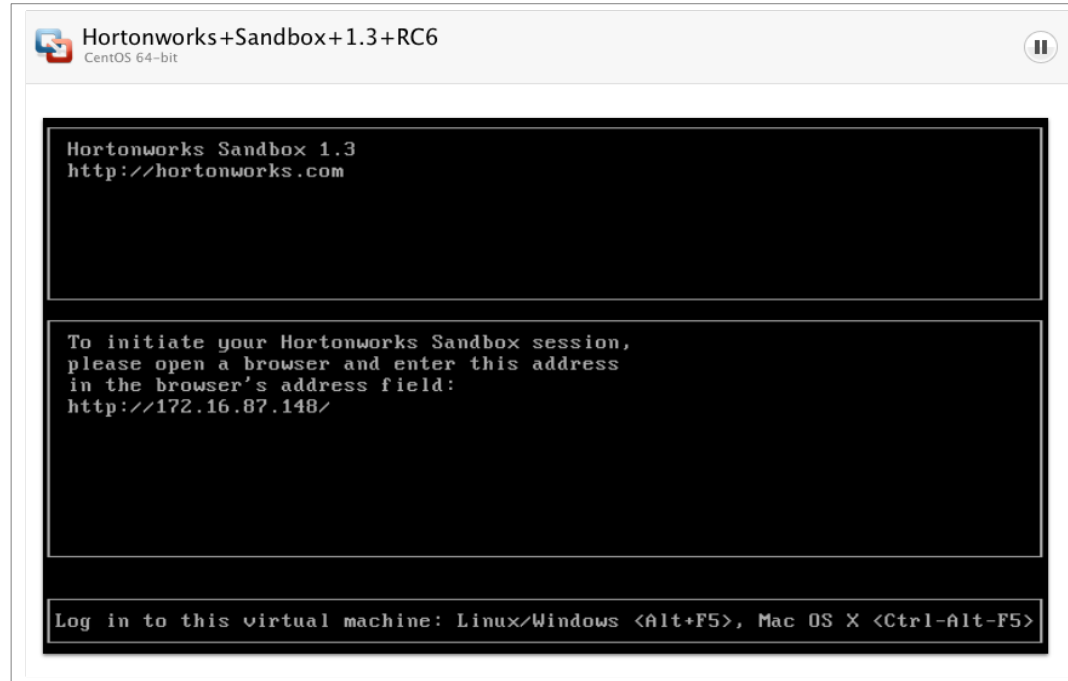
Installing HDP 1.3 Sandbox for VMware

- Configured to use 2 processors
- Uses 2048MB memory
- Network shared with host
sandbox resolves to IP
assigned to VM
- User/password:

root/hadoop
- Listens on ports:

HDFS - sandbox:8020

JobTracker - sandbox:50300



Using HDP 1.3 Sandbox for VMware

- Add to /etc/hosts on your local system (*adjust IP address to the one on startup screen*):

```
172.16.87.148    sandbox
```

- Now you can access Hadoop on the sandbox:

```
$ hadoop dfs -ls hdfs://sandbox:8020/
Found 4 items
drwxr-xr-x   - hdfs    hdfs          0 2013-05-30 13:34 /apps
drwx----- - mapred  hdfs          0 2013-08-23 12:31 /mapred
drwxrwxrwx   - hdfs    hdfs          0 2013-06-10 17:39 /tmp
drwxr-xr-x   - hdfs    hdfs          0 2013-06-10 17:39 /user
```



Hadoop in Pseudo-distributed Mode (Single Node)

- Download Apache Hadoop (hadoop-2.0.6-alpha)
 - <http://hadoop.apache.org/releases.html#Download>
- Create a directory and unzip the download
 - I use ~/Hadoop on my system
- Modify \$HADOOP_INSTALL/etc/hadoop/hadoop-env.sh
 - modify this line : `export JAVA_HOME=${JAVA_HOME}`
 - to be: `export JAVA_HOME="/usr/lib/jvm/java-6-openjdk-amd64"`
or to what your local Java installations home is

Update configuration files in etc/hadoop

core-site.xml

```
<configuration>

  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:8020</value>
    <final>true</final>
  </property>

</configuration>
```

mapred-site.xml

```
<configuration>

  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>

</configuration>
```

hdfs-site.xml

```
<configuration>

  <property>
    <name>dfs.support.append</name>
    <value>true</value>
  </property>

  <property>
    <name>dfs.webhdfs.enabled</name>
    <value>true</value>
  </property>

  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>

</configuration>
```

Update configuration files in etc/hadoop

yarn-site.xml

```
<configuration>

  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce.shuffle</value>
  </property>

  <property>
    <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>
  </property>

</configuration>
```

You can download these config files from:

<https://github.com/trisberg/springone-hadoop/tree/master/hadoop-config/2.0.6-alpha>

Configure your environment settings

hadoop-2.0.6-env

```
export HADOOP_INSTALL=~/.Hadoop/hadoop-2.0.6-alpha
export JAVA_HOME=/usr/lib/jvm/java-6-openjdk-amd64

export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_YARN_HOME=$HADOOP_INSTALL

export HADOOP_CONF_DIR=$HADOOP_INSTALL/etc/hadoop

export PATH=$HADOOP_INSTALL/bin:$HADOOP_INSTALL/sbin:$PATH
```

Configure your SSH settings

Make sure you can ssh to your local system

Create ssh key - *no need to do this if you already have one*

```
ssh-keygen -t dsa -P '' -f ~/.ssh/id_dsa
```

Add the ssh key to authorized keys so you can log in without a password

```
cat ~/.ssh/id_dsa.pub >> ~/.ssh/authorized_keys  
chmod 600 ~/.ssh/authorized_keys
```

Try connecting to local host with ssh (should not be prompted for password)

```
ssh localhost  
exit
```

Let's start by formatting the namenode

```
$ cd ~/Hadoop
$ source hadoop-2.0.6-env
$ hdfs namenode -format
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG:  host = carbon/192.168.0.114
STARTUP_MSG:  args = [-format]
STARTUP_MSG:  version = 2.0.6-alpha
...
Formatting using clusterid: CID-919300bd-2c08-483b-ab8d-a38ce1e31b1c
...
13/08/26 16:15:06 INFO common.Storage: Storage directory /tmp/hadoop-
trisberg/dfs/name has been successfully formatted.
13/08/26 16:15:06 INFO namenode.FSImage: Saving image file /tmp/hadoop-
trisberg/dfs/name/current/fsimage.ckpt_00000000000000000000 using no compression
...
/*****
SHUTDOWN_MSG: Shutting down NameNode at carbon/192.168.0.114
*****/
```

Next, start the Hadoop “cluster”

```
$ start-dfs.sh
13/08/26 16:07:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /home/trisberg/Hadoop/hadoop-2.0.6-
alpha/logs/hadoop-trisberg-namenode-carbon.out
localhost: starting datanode, logging to /home/trisberg/Hadoop/hadoop-2.0.6-
alpha/logs/hadoop-trisberg-datanode-carbon.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/trisberg/Hadoop/hadoop-
2.0.6-alpha/logs/hadoop-trisberg-secondarynamenode-carbon.out

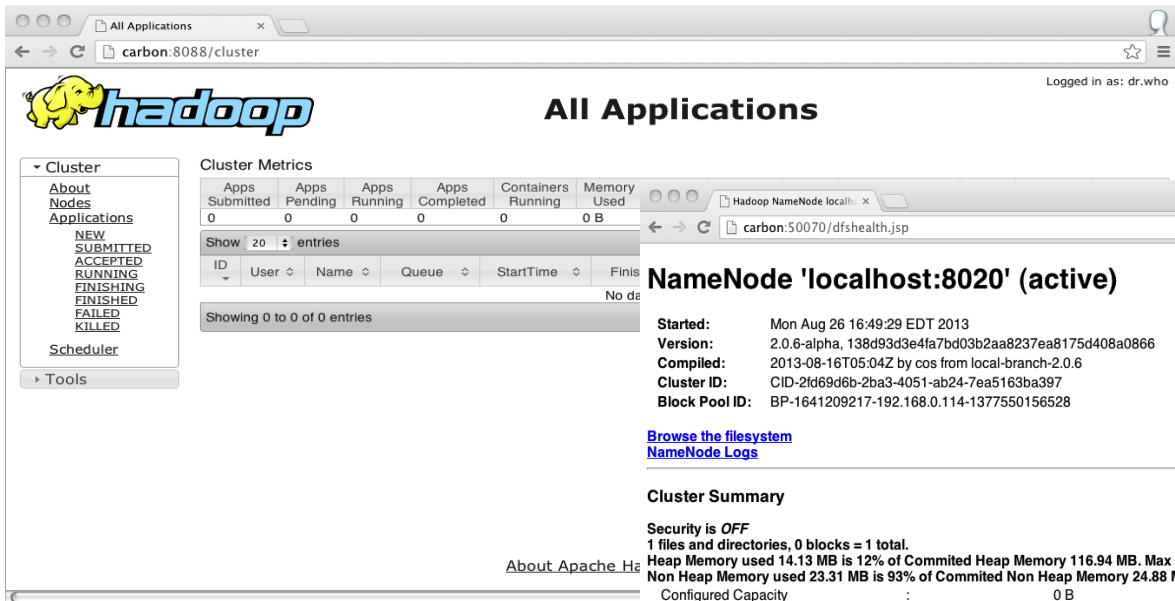
$ start-yarn.sh
starting yarn daemons
starting resourcemanager, logging to /home/trisberg/Hadoop/hadoop-2.0.6-
alpha/logs/yarn-trisberg-resourcemanager-carbon.out
localhost: starting nodemanager, logging to /home/trisberg/Hadoop/hadoop-2.0.6-
alpha/logs/yarn-trisberg-nodemanager-carbon.out
```

Check that all daemons are running

```
$ jps
19995 SecondaryNameNode
19487 NameNode
20183 ResourceManager
19716 DataNode
20591 Jps
20413 NodeManager
```


Check cluster and hdfs web pages

> <http://localhost:8088/>



The screenshot shows the Hadoop All Applications web interface. The top navigation bar includes the Hadoop logo, the title "All Applications", and a login status "Logged in as: dr.who". The left sidebar contains a "Cluster" section with links for "About", "Nodes", "Applications", "NEW SUBMITTED", "ACCEPTED", "RUNNING", "FINISHING", "FINISHED", "FAILED", "KILLED", "Scheduler", and "Tools". The main content area displays "Cluster Metrics" with a table showing 0 apps submitted, pending, running, or completed, and 0 containers running or memory used. Below this is a table for "Show 20 entries" with columns for ID, User, Name, Queue, StartTime, and FinishTime. A message states "Showing 0 to 0 of 0 entries". At the bottom, there is a link for "About Apache Hadoop".

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used
0	0	0	0	0	0 B

ID	User	Name	Queue	StartTime	FinishTime
Showing 0 to 0 of 0 entries					

[About Apache Hadoop](#)

> <http://localhost:50070/>

NameNode 'localhost:8020' (active)

Started: Mon Aug 26 16:49:29 EDT 2013
Version: 2.0.6-alpha, 138d93d3e4fa7bd03b2aa8237ea8175d408a0866
Compiled: 2013-08-16T05:04Z by cos from local-branch-2.0.6
Cluster ID: CID-2fd69d6b-2ba3-4051-ab24-7ea5163ba397
Block Pool ID: BP-1641209217-192.168.0.114-1377550156528

[Browse the filesystem](#)
[NameNode Logs](#)


Cluster Summary

Security is OFF
1 files and directories, 0 blocks = 1 total.
Heap Memory used 14.13 MB is 12% of Committed Heap Memory 116.94 MB. Max Heap Memory is 116.94 MB.
Non Heap Memory used 23.31 MB is 93% of Committed Non Heap Memory 24.88 MB. Max Non Heap Memory is 214 MB.

Configured Capacity	:	0 B		
DFS Used	:	0 B		
Non DFS Used	:	0 B		
DFS Remaining	:	0 B		
DFS Used%	:	100.00%		
DFS Remaining%	:	0.00%		
Block Pool Used	:	0 B		
Block Pool Used%	:	100.00%		
DataNodes usages	:	Min %	Median %	Max %
		0.00%	0.00%	0.00%
				stddev %
				0.00%

For more detail ...

- This has been a brief intro to getting Apache Hadoop installed for development
- Lots more to learn ...



Project Links

- **Source:**
 - <https://github.com/spring-projects/spring-hadoop>
- **Samples:**
 - <https://github.com/spring-projects/spring-hadoop-samples>
- **Project:**
 - <http://projects.spring.io/spring-hadoop/>
- **Forum:**
 - <http://forum.spring.io/forum/spring-projects/data/hadoop>

Learn More. Stay Connected.



We need your feedback -

<http://forum.spring.io/forum/spring-projects/data/hadoop>

- Talk to us on Twitter: @springcentral
- Find Session replays on YouTube: spring.io/video